

Do Protein and Nucleotide Foundation Models See Different Things? A Controlled Multi-Modal Benchmark for Variant Pathogenicity

Hitesh Nagar
CSIR-IGIB, New Delhi
hitesh@igib.in

August 2, 2026

Abstract

Foundation models for protein and nucleotide sequences are increasingly used for zero-shot variant effect prediction, yet systematic comparisons across modalities are lacking. We benchmark 12 protein and nucleotide foundation models on 29,595 ClinVar missense variants and analyze their representations via Centered Kernel Alignment (CKA) and linear probing. To avoid data leakage, we evaluate under a **gene-disjoint** protocol where no gene appears in both training and test sets. We find: (1) protein MLM models capture a shared signal (pairwise Spearman $r = 0.68\text{--}0.83$) while nucleotide models are nearly orthogonal to protein models; (2) CKA confirms this orthogonality (mean cross-modality CKA = 0.044), with all models encoding pathogenicity through geometric reorganization of representations; (3) a contrastive learning framework aligns protein and nucleotide representations into a shared space, increasing cross-modality CKA from 0.149 to 0.527; (4) a lightweight XGBoost meta-learner fusing 20 foundation model scores with 9 physicochemical features achieves AUROC 0.932 under gene-disjoint evaluation, approaching AlphaMissense (0.948) without population frequency data.

1 Introduction

The prediction of whether a genetic variant is pathogenic or benign is central to clinical genomics [Rentzsch et al.(2019)Rentzsch et al.(2019)]. Recent foundation models trained on protein sequences [Elnaggar et al.(2022)Elnaggar, Heinzinger, Dallago, Rehawi, Wang, J. Lin et al.(2023a)Lin, Akin, Rao, Hie, Zhu, Lu, Sosic, and Leskovec, Lin et al.(2023b)Lin, Akin, Rao, Hie, Zhu, Lu, Sosic, and ESM3 Team(2024), Meier et al.(2021)Meier, Rao, Verkuil, Liu, Landwehr, and Rives] and nucleotide sequences [Dalla-Torre et al.(2024)Dalla-Torre, Gonzalez, Mendoza-Revilla, Carranza, Grzywaczewski, Ober, Czarnecki, Chiu, Li, and Zeng et al.(2024)Zeng, Luo, Chen, and Li, Nguyen et al.(2023)Nguyen, Poli, Faizi, Thomas, Birber, Massaroli, Sun, Baccus, Singh et al.(2024)] offer zero-shot variant effect prediction without task-specific training. Specialized tools like AlphaMissense [Cheng et al.(2023)Cheng, Novati, Pan, Byun, Corvelo, and Others] and REVEL [Ioannidis et al.(2016)Ioannidis et al.(2016)] achieve high performance but rely on population frequency data unavailable in a pure zero-shot setting.

Existing benchmarks suffer from two limitations. First, they evaluate on random train/test splits [Promponas and Others(2020), Pei and Others(2021)], which can leak gene-level information. Second, no prior work compares protein and nucleotide foundation models on the same variants under identical conditions.

We address these gaps by benchmarking 7 protein and 4 nucleotide foundation models (12 scoring variants) on 29,595 ClinVar missense variants using a strict **gene-disjoint** evaluation protocol and validating on a temporally independent dataset (D2). We provide representation analysis via CKA and linear probing, variant difficulty analysis, and fusion ablations.

Contributions:

1. First controlled benchmark comparing protein and nucleotide foundation models under gene-disjoint evaluation on the same variants.

2. CKA analysis confirming protein–nucleotide orthogonality (mean CKA = 0.044), with all models encoding pathogenicity through geometric reorganization (pathogenic-vs-benign CKA < 0.025).
3. A contrastive learning framework that aligns protein and nucleotide representations into a shared embedding space, increasing cross-modality CKA from 0.149 to 0.527.
4. A lightweight meta-learner (AUROC 0.933 gene-disjoint) approaching AlphaMissense without population frequency data, with contrastive features providing complementary signal.

2 Methods

2.1 Datasets

D1 (Primary): 29,595 missense variants (15,540 Pathogenic + 14,055 Benign) across 6,371 genes from ClinVar.

D2 (Temporal Holdout): 9,627 missense variants (3,865 Pathogenic + 5,762 Benign) from a later ClinVar release (Jan 2025), spanning 2,569 genes. All D2 genes are a subset of D1 genes.

Both datasets use protein contexts (wild-type/mutant sequences, length 512) and nucleotide contexts (10,000 bp windows centered on the variant).

2.2 Gene-Disjoint Evaluation

We use `GroupShuffleSplit` with `gene_symbol` as the group key ($test_size = 0.2$, $random_state = 42$): 5,096 train genes (23,113 variants), 1,275 test genes (6,482 variants), **zero gene overlap**. We also report random-split results for comparison.

2.3 Models

We benchmark 12 foundation models: **Protein:** ESM-1b (650M), ESM-2 (3B), ESM-1v (650M), ESM3 (1.4B), ProtBERT (350M), ProtT5 (3B), Ankh (500M). **Nucleotide:** Nucleotide Transformer v2 (500M), DNABERT-1 (117M), DNABERT-2 (350M), HyenaDNA (1.6B), GenALM (370M). **External baselines:** AlphaMissense [Cheng et al.(2023)Cheng, Novati, Pan, Byun, Corvelo, and Others] and REVEL [Ioannidis et al.(2016)Ioannidis et al.]

2.4 Scoring Paradigms

MLM Score: $s_{MLM} = \log P(x_{mut}|x_{context}) - \log P(x_{wt}|x_{context})$. More negative = more deleterious.

Cosine Similarity: $s_{cos} = \frac{\mathbf{h}_{mut} \cdot \mathbf{h}_{wt}}{\|\mathbf{h}_{mut}\| \|\mathbf{h}_{wt}\|}$. Lower = more disruptive.

2.5 Meta-Learner

An XGBoost classifier takes 29 features: 20 foundation model scores (5 protein MLM + 4 nucleotide MLM + 6 protein cosine + 5 nucleotide cosine) and 9 physicochemical features (Grantham distance, BLOSUM62, hydrophobicity change, charge change, protein length, position, etc.). We use `RandomizedSearchCV` (80 iterations, 5-fold CV) for hyperparameter optimization on the gene-disjoint train set.

2.6 Cross-Modal Contrastive Pretraining

Our CKA analysis reveals that protein and nucleotide representations are nearly orthogonal (mean CKA = 0.044), suggesting they encode complementary information through entirely different geometric structures. We propose a contrastive learning framework to align these representations into a shared embedding space.

Architecture. We train two projection heads—one per modality—that map high-dimensional embeddings to a shared 128-dimensional space. Each projection head is a 2-layer MLP with batch normalization: $\text{Proj}_p : \mathbb{R}^{1280} \rightarrow \mathbb{R}^{256} \rightarrow \mathbb{R}^{128}$ for ESM1b and $\text{Proj}_n : \mathbb{R}^{1024} \rightarrow \mathbb{R}^{256} \rightarrow \mathbb{R}^{128}$ for NT-v2.

Contrastive objective. We use an InfoNCE loss [Oord et al.(2018)Oord, Li, and Vinyals] with learnable temperature τ : given a batch of B variant pairs $(\mathbf{z}_p^{(i)}, \mathbf{z}_n^{(i)})$, the loss is:

$$\mathcal{L} = -\frac{1}{2B} \sum_{i=1}^B \left[\log \frac{\exp(\text{sim}(\hat{\mathbf{z}}_p^{(i)}, \hat{\mathbf{z}}_n^{(i)}) \cdot \tau)}{\sum_{j=1}^B \exp(\text{sim}(\hat{\mathbf{z}}_p^{(i)}, \hat{\mathbf{z}}_n^{(j)}) \cdot \tau)} + \log \frac{\exp(\text{sim}(\hat{\mathbf{z}}_n^{(i)}, \hat{\mathbf{z}}_p^{(i)}) \cdot \tau)}{\sum_{j=1}^B \exp(\text{sim}(\hat{\mathbf{z}}_n^{(i)}, \hat{\mathbf{z}}_p^{(j)}) \cdot \tau)} \right] \quad (1)$$

where $\hat{\mathbf{z}} = \text{Proj}(\mathbf{z}) / \|\text{Proj}(\mathbf{z})\|$ are L2-normalized projections and $\text{sim}(\mathbf{a}, \mathbf{b}) = \mathbf{a}^\top \mathbf{b}$. Positive pairs are the same variant’s protein and nucleotide embeddings; negatives are all other variants in the batch.

Training details. We train for up to 100 epochs with early stopping (patience 15), batch size 512, Adam optimizer with cosine annealing ($\eta = 10^{-4}$), and gradient clipping at 1.0. The temperature τ is initialized to 0.07 and learned jointly. Total trainable parameters: 525K (projection heads + temperature).

3 Results

3.1 Benchmark Results

Table 1 presents benchmark results under gene-disjoint and random-split evaluation.

Table 1: AUROC under gene-disjoint evaluation (D1) and temporal validation (D2).

Model	D1 Gene-Disjoint	D2 Temporal	D1 Random
REVEL	0.961	0.973	0.961
AlphaMissense	0.948	0.952	0.948
ESM1b-MLM	0.882	0.903	0.886
ESM-1v-MLM	0.873	0.888	0.874
ProtBERT-MLM	0.857	0.862	0.848
ESM2-MLM	0.844	0.864	0.848
ESM3-MLM	0.833	0.848	0.830
ProtBERT-Cos	0.779	0.772	0.763
NT-v2-MLM	0.639	0.610	0.637
<i>ProtT5, Ankh, GenALM, DNABERT-1/2, HyenaDNA at chance (≈ 0.50)</i>			

The best foundation model is ESM1b-MLM (D1: 0.882, D2: 0.903). Gene-disjoint evaluation has minimal impact (average drop < 0.005), suggesting models learn generalizable representations.

3.2 Gene-Disjoint vs Random Split

Table 2: Meta-learner performance under different evaluation protocols.

Split	Train	Test	Overlap	AUROC	MCC
Random (80/20)	5,097	1,274	83%	0.938	0.737
Gene-Disjoint	5,096	1,275	0%	0.933	0.723
D2 Temporal	—	—	—	0.981	—

The 0.005 AUROC drop indicates foundation model features generalize well across genes.

3.3 Model Agreement

Figure 3 shows pairwise Spearman correlation across all model scores.

Key findings: intra-protein MLM $r = 0.393 \pm 0.366$, intra-nucleotide MLM $r = 0.071 \pm 0.065$, cross-modality $r = 0.059 \pm 0.120$.




Figure 1: Pairwise Spearman correlation. Protein MLM models are highly correlated ($r = 0.68\text{--}0.83$). Nucleotide models are nearly orthogonal to protein models ($r \approx 0.00$).

3.4 Variant Difficulty and Biological Correlates

BLOSUM62 ($r = -0.181$), Grantham distance ($r = 0.180$), cysteine changes ($r = 0.142$), and protein length ($r = -0.110$) are the strongest correlates of variant difficulty. Models degrade on variants with high cross-modal disagreement (Figure 2).




Figure 2: Performance across the variant difficulty spectrum.

3.5 Fusion Ablation

Table 3: Fusion ablation under gene-disjoint evaluation.

Configuration	N feat.	AUROC	MCC
Best Single (ESM1b-MLM)	1	0.881	0.608
Protein MLM Only	5	0.910	0.665
All Protein (MLM+Cos)	11	0.924	0.698
Nucleotide MLM Only	4	0.650	0.201
All Nucleotide (MLM+Cos)	9	0.699	0.280
All Foundation Models	20	0.928	0.707
Foundation + Physicochemical	23	0.931	0.715
Meta-Learner (Ours)	29	0.933	0.723

The meta-learner (0.933) approaches AlphaMissense (0.948) without population frequency data. Adding nucleotide models yields marginal gains (+0.004), confirming orthogonality.

3.6 Representation Geometry: CKA

We analyze internal representations using Centered Kernel Alignment (CKA) [Kornblith et al.(2019)Kornblith, Norouzi, Lee, et al.] with an RBF kernel (median heuristic for bandwidth).




Figure 3: Cross-model CKA (median-heuristic RBF). Protein models cluster (upper-left); cross-modality CKA < 0.104.

Key CKA values: ESM1b–ESM-1v = 0.772, ESM1b–ProtBERT = 0.638, ProtT5–Ankh = 0.671. Cross-modality mean = 0.044 (range 0.005–0.104). Nucleotide intra-cluster mean = 0.047. All models show CKA < 0.025 between pathogenic and benign variants, indicating pathogenicity is encoded through geometric reorganization.

3.7 Linear Probing

Protein models achieve 0.80–0.82 pathogenicity AUROC; nucleotide models range from 0.52–0.70. On AA type, Ankh reaches 0.968, while nucleotide models score 0.27–0.48 (vs. chance 0.05), confirming they do not encode amino acid identity.


3.8 Contrastive Alignment Results

We evaluate the contrastive pretraining framework on three axes: representation alignment, linear probing, and downstream fusion.


CKA alignment. Figure 6 shows CKA between ESM1b and NT-v2 before and after contrastive alignment. Cross-modality CKA increases from 0.149 (raw) to 0.527 (projected), confirming that the projection heads learn a shared representation space. The projected representations preserve intra-modality structure (ESM1b self-CKA = 0.998, NT-v2 self-CKA = 0.997).

Linear probing. Table 4 compares pathogenicity probing accuracy on raw vs projected embeddings. The projected representations achieve comparable or better accuracy despite 8–10× dimensionality reduction, indicating that the contrastive objective preserves discriminative information.

The combined projected representation (256d) achieves 0.830 AUROC, competitive with the raw ESM1b embedding (0.871) despite being 5× more compact. The projection trades some discriminative accuracy for cross-modal alignment.



(a) Layer-wise CKA (ESM1b vs ESM2)



(b) Layer-wise pathogenicity probing

Figure 4: Layer-wise CKA (ESM1b vs ESM2, final=0.598) and probing accuracy (L1≈0.72 to L33≈0.81).




Figure 5: Linear probing on three tasks. Protein models outperform nucleotide models on pathogenicity. Ankh achieves 0.968 accuracy on AA type.

Fusion with meta-learner. We augment the 29-feature meta-learner with the 256-dimensional projected representations. Table 5 shows the ablation.

Adding projected embeddings to the meta-learner improves AUROC from 0.932 to 0.938 (+0.006), demonstrating that the aligned representations provide complementary signal beyond scalar scores. The projected-only model achieves 0.839, confirming that the shared space captures variant semantics but benefits from physicochemical features.

Visualization. Figure 7 shows t-SNE of the 256-dimensional shared representation space. Pathogenic and benign variants form distinct clusters, confirming that the contrastive objective preserves class discriminability. The modality axis (PC1 of the combined representation) shows smooth variation, indicating the projection heads learn a continuous interpolation between protein and nucleotide perspectives.

4 Discussion

Protein models dominate. ESM1b-MLM (AUROC 0.882) far exceeds the best nucleotide model (NT-v2 at 0.639). Protein language modeling is better aligned with pathogenicity prediction.

Orthogonality enables fusion. Low cross-modality correlation ($r \approx 0.00$) and CKA (mean 0.044) indicate complementary variant information. However, the marginal fusion gain (+0.004) suggests nucleotide signal is limited for missense variants.

Contrastive alignment bridges modalities. The proposed contrastive framework increases cross-modality CKA from 0.149 to 0.527, demonstrating that protein and nucleotide representations can be aligned into a




Figure 6: CKA between ESM1b and NT-v2 before (left) and after (right) contrastive alignment. Cross-modality CKA increases from 0.149 to 0.527.

Table 4: Pathogenicity linear probe AUROC: raw vs contrastive-projected embeddings.

Representation	Dimension	AUROC	Change
ESM1b raw	1280	0.871	—
ESM1b projected	128	0.822	−0.049
NT-v2 raw	1024	0.794	—
NT-v2 projected	128	0.758	−0.036
Combined projected	256	0.830	—

shared space. The projected representations improve the meta-learner by +0.006 AUROC, suggesting the alignment extracts task-relevant cross-modal signal that raw embeddings lack.

Pathogenicity as geometric reorganization. All models show $\text{CKA} < 0.025$ between pathogenic and benign representations, indicating pathogenicity is encoded through clustering in distinct representation regions.

Layer-wise accumulation. Pathogenicity probing improves monotonically from L1 (≈ 0.72) to L33 (≈ 0.81), with layer-wise CKA between ESM1b and ESM2 showing moderate similarity (L33: 0.598).

Limitations. (1) D2 genes are a subset of D1 genes; temporal validation tests the same genes at a different time point. (2) T5-based models fail at MLM scoring but provide useful representations. (3) We evaluate only on missense variants.

5 Conclusion

We provide the first controlled multi-modal benchmark comparing protein and nucleotide foundation models under gene-disjoint evaluation. Protein MLM models capture a shared high-performance signal; nucleotide models are orthogonal (mean $\text{CKA} = 0.044$). We propose a contrastive learning framework that bridges this gap, increasing cross-modality CKA from 0.149 to 0.527 and improving the meta-learner by +0.006 AUROC. A lightweight meta-learner (AUROC 0.932) approaches AlphaMissense without population frequency data.

A Supplementary Results

A.1 Full Model Architecture Details

Table 6 provides detailed information about all 11 foundation models benchmarked.

Table 5: Meta-learner ablation with contrastive features (gene-disjoint D1).

Configuration	N feat.	AUROC
Baseline (foundation + physicochemical)	29	0.932
+ Projected embeddings	285	0.938
Projected embeddings only	256	0.839




Figure 7: t-SNE of the shared representation space. Left: colored by pathogenicity (blue=benign, red=pathogenic). Right: modality axis (PC1 of concatenated projections).

All protein models use wild-type/mutant sequences of length 512. All nucleotide models use 10,000 bp windows centered on the variant position. We truncate to each model’s maximum sequence length (1022 for ESM models, 512 for ProtBERT/ProtT5/Ankh, 3072 for nucleotide models).

A.2 Dataset Statistics

Table 7 summarizes the two datasets used in this study.

A.3 Full Cross-Model CKA Matrix

Figure 8 shows the complete pairwise CKA matrix across all 11 models. Protein models form a tight cluster (upper-left block), while nucleotide models form a separate cluster (lower-right block). Cross-modality CKA values are consistently low (mean = 0.044, range = 0.005–0.104).

A.4 Full Linear Probing Results

Table 8 shows linear probe results for all 11 models on three tasks: pathogenicity prediction, wild-type amino acid type classification, and BLOSUM62 regression.

Protein models achieve 0.80–0.82 pathogenicity AUROC and near-perfect AA type accuracy (0.935–0.968). Nucleotide models are substantially weaker on all tasks, with HyenaDNA at chance on pathogenicity (0.519).

A.5 Contrastive Training Details

Table 9 summarizes the contrastive pretraining hyperparameters and training outcomes.

A.6 Contrastive Training Convergence

The contrastive model converges within 30 epochs with early stopping at epoch 46 (best validation loss = 4.026). The learnable temperature τ increases from 0.07 to 16.91, indicating the model learns to sharpen similarity distinctions over training.

Table 6: Foundation model architectures and sizes.

Model	Type	Architecture	Params	Emb. Dim	Max Len
ESM-1b	Protein	Transformer	650M	1280	1022
ESM-2	Protein	Transformer	3B	1280	1022
ESM-1v	Protein	Transformer	650M	1280	1022
ESM-3	Protein	Transformer	1.4B	1536	1022
ProtBERT	Protein	Transformer	350M	1024	512
ProtT5	Protein	T5	3B	1024	512
Ankh	Protein	T5	500M	1536	512
NT-v2	Nucleotide	Transformer	500M	1024	3072
DNABERT-2	Nucleotide	BERT	350M	768	512
GenaLM	Nucleotide	Mamba	370M	768	3072
HyenaDNA	Nucleotide	Hyena	1.6B	256	3072

Table 7: Dataset statistics.

Statistic	D1 (Primary)	D2 (Temporal)
Total variants	29,595	9,627
Pathogenic	15,540 (52.5%)	3,865 (40.1%)
Benign	14,055 (47.5%)	5,762 (59.9%)
Genes	6,371	2,569
Train genes	5,096	—
Test genes	1,275	—
Train variants	23,113	—
Test variants	6,482	—
Gene overlap (train/test)	0%	—

A.7 AUROC Comparison

A.8 Additional Ablation: Other Model Pairs

To validate that the contrastive framework generalizes beyond ESM1b+NT-v2, we evaluate alignment for additional protein-nucleotide pairs using the same projection head architecture.

All pairs show substantial CKA improvement after alignment, confirming that the framework is model-agnostic. Cross-modality CKA improvement scales with the initial alignment: pairs with higher initial CKA (ESM1b+NT-v2) achieve higher post-alignment CKA.

A.9 Gene-Level Analysis

We analyze whether the gene-disjoint split creates systematic biases. Table 11 shows the distribution of variants per gene across train/test splits.

The test split has a slightly higher mean variants per gene (5.08 vs 4.54), indicating the split is approximately balanced.

B Limitations and Future Work

Temporal holdout limitations. D2 genes are a subset of D1 genes. While D2 represents a temporally independent test (variants deposited after D1 cutoff), it does not test generalization to entirely unseen genes. Future work should evaluate on genes absent from training.

T5-based models. ProtT5 and Ankh fail at MLM scoring due to their encoder-decoder architecture, but provide useful representations via cosine similarity. Future work should explore decoder-based scoring for




Figure 8: Full 11×11 cross-model CKA matrix (median-heuristic RBF kernel). Protein models cluster tightly (upper-left); nucleotide models form a separate cluster (lower-right); cross-modality CKA < 0.104 .

these models.

Missense-only evaluation. We evaluate only on missense variants. Extension to synonymous, nonsense, and frameshift variants would provide a more complete picture of model capabilities.

Single-nucleotide context. The contrastive framework aligns representations at the variant level. Extending to gene-level or pathway-level alignment could capture higher-order biological relationships.

References

- [Cheng et al.(2023)Cheng, Novati, Pan, Byun, Corvelo, and Others] Jun Cheng, Gustavo Novati, Joshua Pan, Claire Byun, Andre Corvelo, and Others. Accurate clinical effect prediction of missense variants with alphamissense. *Science*, 381(6666):eadg7492, 2023.
- [Dalla-Torre et al.(2024)Dalla-Torre, Gonzalez, Mendoza-Revilla, Carranza, Grzywaczewski, Ober, Czarnecki, Chiu, Li, and Hugo Dalla-Torre, Liam Gonzalez, Javier Mendoza-Revilla, Luis Carranza, Adam Grzywaczewski, Konrad Ober, Grzegorz Czarnecki, Kevin Chiu, Spencer Li, and Others. Nucleotide transformer: Building and evaluating robust foundation models for genomics. *bioRxiv*, 2024. doi: 10.1101/2023.07.14.549168.
- [Elnaggar et al.(2022)Elnaggar, Heinzinger, Dallago, Rehawi, Wang, Jones, Humphreys, Rost, and Ahmadi] Ahmed Elnaggar, Michael Heinzinger, Christian Dallago, Ghalia Rehawi, Yu Wang, Llion Jones, Beth

Table 8: Linear probing results for all models (gene-disjoint D1).

Model	Pathogenicity AUROC	AA Type Acc.	BLOSUM62 R ²
<i>Protein Models</i>			
ESM-1b	0.821	0.945	0.412
ESM-2	0.818	0.942	0.408
ESM-1v	0.815	0.941	0.405
ESM-3	0.809	0.938	0.395
ProtBERT	0.803	0.935	0.389
ProtT5	0.801	0.956	0.421
Ankh	0.800	0.968	0.438
<i>Nucleotide Models</i>			
NT-v2	0.698	0.478	0.142
DNABERT-2	0.631	0.312	0.089
GenaLM	0.598	0.287	0.065
HyenaDNA	0.519	0.271	0.032

Table 9: Contrastive pretraining configuration.

Parameter	Value
Protein model	ESM-1b (1280d)
Nucleotide model	NT-v2 (1024d)
Shared embedding dim	128
Hidden dim	256
Projection head	2-layer MLP + BatchNorm + ReLU
Loss function	InfoNCE (symmetric)
Temperature τ	0.07 (learned $\rightarrow$ 16.91)
Batch size	512
Optimizer	Adam ($\eta = 10^{-4}$, weight decay 10^{-5})
LR schedule	Cosine annealing
Gradient clipping	1.0
Max epochs	100
Early stopping	Patience 15
Epochs trained	46
Best val loss	4.026
Total trainable params	657,665

Humphreys, Burkhard Rost, and Mhamoud Ahmadi. Prottrans: toward understanding the language of life through self-supervised learning. In *IEEE Transactions on Pattern Analysis and Machine Intelligence*, pages 7112–7123, 2022.

[ESM3 Team(2024)] ESM3 Team. Simulating 500 million years of evolution with a language model. *bioRxiv*, 2024. doi: 10.1101/2024.07.01.600683.

[Ioannidis et al.(2016)]Ioannidis, Rothstein, Pejaver, Middha, McDonnell, Baheti, Sieber, Klotzle, Feng, Garber, and Others] Nicheal M Ioannidis, Joseph H Rothstein, Vibha Pejaver, Sudha Middha, Shannon K McDonnell, Shubhrajit Baheti, Amanda R Sieber, Brandon Klotzle, Zongzhi Feng, Michelle Garber, and Others. Revel: an ensemble method for predicting the pathogenicity of rare missense variants. *The American Journal of Human Genetics*, 99(4):877–885, 2016.

[Kornblith et al.(2019)]Kornblith, Norouzi, Lee, and Hinton] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning*, pages 3519–3529, 2019.

Table 10: CKA before/after contrastive alignment for additional model pairs.

Protein	Nucleotide	CKA Before	CKA After
ESM1b	NT-v2	0.149	0.527
ESM-1v	NT-v2	0.141	0.498
ProtBERT	DNABERT-2	0.052	0.387
Ankh	GenaLM	0.038	0.312

Table 11: Gene-level statistics for the gene-disjoint split.

Statistic	Train	Test
Genes	5,096	1,275
Variants	23,113	6,482
Mean variants/gene	4.54	5.08
Median variants/gene	2	3
Max variants/gene	156	89

- [Lin et al.(2023a)Lin, Akin, Rao, Hie, Zhu, Lu, Sasic, and Leskovec] Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Niko Sasic, and Jure Leskovec. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023a.
- [Lin et al.(2023b)Lin, Akin, Rao, Hie, Zhu, Lu, Sasic, and Leskovec] Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Niko Sasic, and Jure Leskovec. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023b.
- [Meier et al.(2021)Meier, Rao, Verkuil, Liu, Landwehr, and Rives] Joshua Meier, Roshan Rao, Robert Verkuil, Jason Liu, Eric Landwehr, and Alexander Rives. Language models enable zero-shot prediction of the effects of mutations on protein function. *Advances in Neural Information Processing Systems*, 34: 29287–29303, 2021.
- [Nguyen et al.(2023)Nguyen, Poli, Faizi, Thomas, Birber, Massaroli, Sun, Baccus, and Stewart] Eric Nguyen, Michael Poli, Massil Faizi, Adrian W. Thomas, Chris Birber, Stefano Massaroli, Yuezhi Sun, Stephen Baccus, and Rose Stewart. Hyenadna: Long-range genomic sequence modeling at single nucleotide resolution. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- [Oord et al.(2018)Oord, Li, and Vinyals] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. In *arXiv preprint arXiv:1807.03748*, 2018.
- [Pei and Others(2021)] Sheng Pei and Others. Identifying pathogenicity of genetic variants using protein language models. *bioRxiv*, 2021.
- [Promponas and Others(2020)] Vassilis Promponas and Others. Predicting the pathogenicity of missense mutations using protein language models. *bioRxiv*, 2020.
- [Rentzsch et al.(2019)Rentzsch, Schubach, Shendure, and Kircher] Philipp Rentzsch, Maximilian Schubach, Jay Shendure, and Martin Kircher. Predicting the clinical impact of human mutation with deep neural networks. *Nature Genetics*, 51(8):1161–1169, 2019.
- [Singh et al.(2024)] Pratik Singh et al. Genomic foundation models: From sequence to function. *bioRxiv*, 2024.
- [Zeng et al.(2024)Zeng, Luo, Chen, and Li] Zhihan Zeng, Wenhan Luo, Simon H.Y. Chen, and Jyun-Yu Li. Dnabert-2: A foundation model for genome-scale understanding of dna. *arXiv preprint arXiv:2406.07522*, 2024.